

Your settings for this Block Poster are:

Pages Wide
Orientation
Paper Format
Border Setting

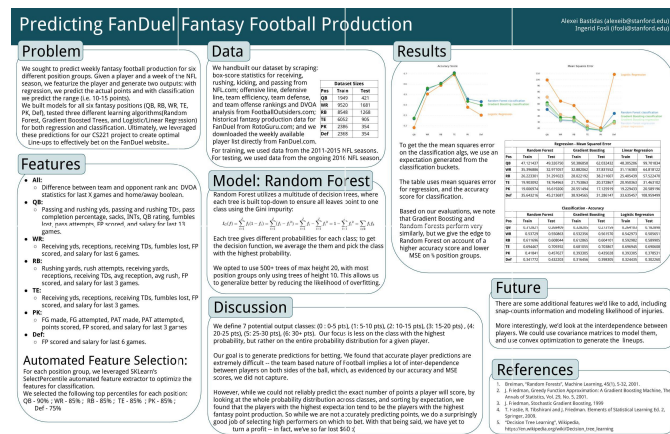
4

PORTRAIT
LETTER

Top: 8, Right: 8, Bottom: 9, Left: 9

You can find tips on printing, assembling and putting up your poster at www.blockposters.com

Enjoy your Block Poster!



Predicting FanDuel

Problem

We sought to predict weekly fantasy football production different position groups. Given a player and a week of the season, we featurize the player and generate two outputs. With regression, we predict the actual points and with classification we predict the range (i.e. 10-15 points).

We built models for all six fantasy positions (QB, RB, WR, PK, Def), tested three different learning algorithms (Random Forest, Gradient Boosted Trees, and Logistic/Linear Regression) for both regression and classification. Ultimately, we leveraged these predictions for our CS221 project to create optimal Line-ups to effectively bet on the FanDuel website..

Features

- **All:**
 - Difference between team and opponent rank and statistics for last X games and home/away boolean
- **QB:**
 - Passing and rushing yds, passing and rushing TDs, completion percentage, sacks, INTs, QB rating, fumbles lost, pass attempts, FP scored, and salary for last

Fantasy Football Pro

Data

We handbuilt our dataset by scraping: box-score statistics for receiving, rushing, kicking, and passing from NFL.com; offensive line, defensive line, team efficiency, team defense, and team offense rankings and DVOA analysis from FootballOutsiders.com; historical fantasy production data for FanDuel from RotoGuru.com; and we downloaded the weekly available player list directly from FanDuel.com.

For training, we used data from the 2011-2015 NFL season. For testing, we used data from the ongoing 2016 season.

Data	
Pos	Trails
QB	194
WR	952
RB	854
TE	605
PK	238
Def	236

Model: Random Forest

Random Forest utilizes a multitude of decision trees. Each tree is built top-down to ensure all leaves belong to a single class using the Gini impurity:

$$I_G(f) = \sum_i f_i(1 - f_i) = \sum_i (f_i - f_i^2) = \sum_i f_i - \sum_i f_i^2 = 1 - \sum_i f_i^2$$

roduction

Results

set Sizes

in	Test
9	421
0	1681
8	1268
2	905
6	354
8	354

JFL seasons.
NFL season.

rest

rees, where
point to one

$$f_i^2 = \sum f_i f_k$$



To get the the mean squares error on the classification algs, we use an expectation generated from the classification buckets.

The table uses mean squares error for regression, and the accuracy score for classification.

Based on our evaluations, we note that Gradient Boosting and Random Forests perform very

Pos
QB
WR
RB
TE
PK
Def
Pos

Alexei Bastidas (alexeib@stanford.edu)
Ingerid Fosli (ifosli@stanford.edu)



Regression - Mean Squared Error

Random Forest		Gradient Boosting		Linear Regression	
Train	Test	Train	Test	Train	Test
47.121437	49.326750	50.386858	62.032432	48.305206	99.701834
35.396886	32.971057	32.882062	37.831552	31.116383	64.818122
26.223301	31.291023	28.022192	38.211607	25.405439	57.522470
19.903092	18.764963	21.753863	20.372867	20.950363	31.463102
19.006974	16.619200	20.551494	17.125919	19.229433	20.589196
35.643216	45.213687	38.934565	31.286147	33.635457	108.959499

Classification - Accuracy

Random Forest		Gradient Boosting		Logistic Regression	
Train	Test	Train	Test	Train	Test

receiving yds, receiving TDs, receiving TDs, receiving TDs, and salary for last 6 games.

- **WR:**
 - Receiving yds, receptions, receiving TDs, fumbles scored, and salary for last 6 games.
- **RB:**
 - Rushing yards, rush attempts, receiving yards, receptions, receiving TDs, avg reception, avg rush scored, and salary for last 3 games.
- **TE:**
 - Receiving yds, receptions, receiving TDs, fumbles scored, and salary for last 3 games.
- **PK:**
 - FG made, FG attempted, PAT made, PAT attempted, points scored, FP scored, and salary for last 3 games.
- **Def:**
 - FP scored and salary for last 6 games.

Automated Feature Selection

For each position group, we leveraged SKLearn's SelectPercentile automated feature extractor to optimize features for classification.

We selected the following top percentiles for each position group:
QB - 90% ; WR - 85% ; RB - 85% ; TE - 85% ; PK - 85% ;
Def - 75%

Each tree gives different probabilities for each class. In the decision function, we average the probabilities and choose the class with the highest probability.

We opted to use 500+ trees of max height 20, with position groups only using trees of height 10. This helps to generalize better by reducing the likelihood of overfitting.

Discussion

We define 7 potential output classes: (0 : 0-5 pts), (1: 5-10 pts), (2: 10-15 pts), (3: 15-20 pts), (4: 20-25 pts), (5: 25-30 pts), (6: 30+ pts). Our focus is on the probability, but rather on the entire probability distribution.

Our goal is to generate predictions for betting. This is extremely difficult -- the team based nature of football, the interaction between players on both sides of the ball, which we did not capture.

However, while we could not reliably predict the outcome, looking at the whole probability distribution across all players, we found that the players with the highest expected fantasy point production. So while we are not a good job of selecting high performers on which to bet, we turn a profit -- in fact, we've so far lost \$6

QB
WR
RB
TE
PK
Def

$\frac{1}{n} \sum_{i \neq k}$
 class; to get
 pick the class

with most
 this allows us
 of overfitting.

Random Forests perform very
 similarly, but we give the edge to
 Random Forest on account of a
 higher accuracy score and lower
 MSE on 5% position groups.

Fu

There
snap-c

More
player
and u

s), (1: 5-10 pts), (2: 10-15 pts), (3: 15-20 pts) , (4:
 s is less on the class with the highest
 distribution for a given player.

We found that accurate player predictions are
 football implies a lot of inter-dependence
 , as evidenced by our accuracy and MSE

Re

1. Bre
2. J. Fr
Anr
3. J. Fr
4. T. F
Spr
5. "De
http

e exact number of points a player will score, by
 cross classes, and sorting by expectation, we
 tion tend to be the players with the highest
 ccurately predicting points, we do a surprisingly
 to bet. With that being said, we have yet to
 50 :(

0.312821	0.268409	0.328205	0.273159	0.264103	0.182898
0.53729	0.550863	0.532356	0.561570	0.542973	0.505651
0.611696	0.608044	0.612865	0.604101	0.592982	0.589905
0.694467	0.709392	0.681055	0.703867	0.696945	0.690608
0.41841	0.457627	0.393305	0.435028	0.393305	0.378531
0.341772	0.432203	0.316456	0.398305	0.324435	0.302260

ture

are some additional features we'd like to add, including counts information and modeling likelihood of injuries.

interestingly, we'd look at the interdependence between s. We could use covariance matrices to model them, se convex optimization to generate the lineups.

ferences

iman, "Random Forests", Machine Learning, 45(1), 5-32, 2001.

riedman, Greedy Function Approximation: A Gradient Boosting Machine, The als of Statistics, Vol. 29, No. 5, 2001.

riedman, Stochastic Gradient Boosting, 1999

lastie, R. Tibshirani and J. Friedman. Elements of Statistical Learning Ed. 2, inger, 2009.

cision Tree Learning", Wikipedia,

os://en.wikipedia.org/wiki/Decision_tree_learning

